


How ML works, and its application

Zhaohui Cheng
2020/04/21

Outcome

- Is machine learning magic?
- Does it require complex maths?
- How to read this?
- What is its application in our field?



How ML works, e.g. image classification



(assume given set of discrete labels)
{dog, cat, truck, plane, ...}



cat



08	02	22	97	38	15	00	40	00	75	04	05	07	78	52	12	50	97	17	68
49	49	99	40	17	81	18	57	60	87	17	40	96	43	05	11	04	56	62	00
01	49	31	73	55	79	14	29	93	71	90	82	01	33	30	03	49	13	36	65
52	70	95	23	04	60	11	62	85	47	95	94	01	32	54	71	37	02	36	91
22	31	16	71	51	62	07	59	41	92	36	54	22	40	40	28	66	33	13	80
24	47	54	00	99	05	45	02	44	75	33	53	78	36	84	20	35	17	12	50
12	98	81	28	64	23	47	10	26	38	40	47	59	54	70	66	18	38	64	70
67	26	20	68	02	62	12	20	95	63	94	39	43	08	40	91	66	49	94	21
24	55	58	05	66	73	99	26	97	17	78	78	96	83	14	88	34	89	63	72
21	36	23	09	75	00	76	44	20	45	35	14	00	61	33	97	34	31	33	95
78	17	53	28	22	75	31	67	15	94	03	80	04	62	16	14	09	53	56	92
16	39	05	42	96	35	31	47	55	58	88	24	00	17	54	24	36	29	85	57
96	56	00	48	35	71	89	07	05	44	44	37	44	60	21	58	51	54	17	58
19	80	81	48	05	34	47	69	29	73	92	13	86	52	17	77	04	89	55	40
04	52	08	83	97	35	98	14	07	97	57	32	14	24	24	79	33	27	88	66
79	46	68	87	57	62	28	72	03	46	35	67	46	55	12	32	43	93	53	49
04	42	14	73	15	85	33	11	24	94	72	18	08	46	29	32	40	42	76	36
20	69	36	41	72	30	23	88	14	17	85	69	82	67	59	85	74	04	36	16
20	73	35	29	78	31	90	01	74	31	49	72	48	18	61	16	23	57	05	54
01	70	54	71	83	51	54	69	16	92	33	45	61	43	52	01	87	27	67	48

What the computer sees

An example dataset: CIFAR-10

Canadian Institute For Advanced Research

Collected by Alex Krizhevsky,
Vinod Nair, and Geoffrey Hinton.

60,000 colour images
Each 32x32x3 pixels

10 classes,
6000 images/class

50,000 training images
10,000 test images

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck





[32x32x3]

array of numbers 0...1
(3072 numbers total)

image parameters

$f(\mathbf{x}, \mathbf{W})$

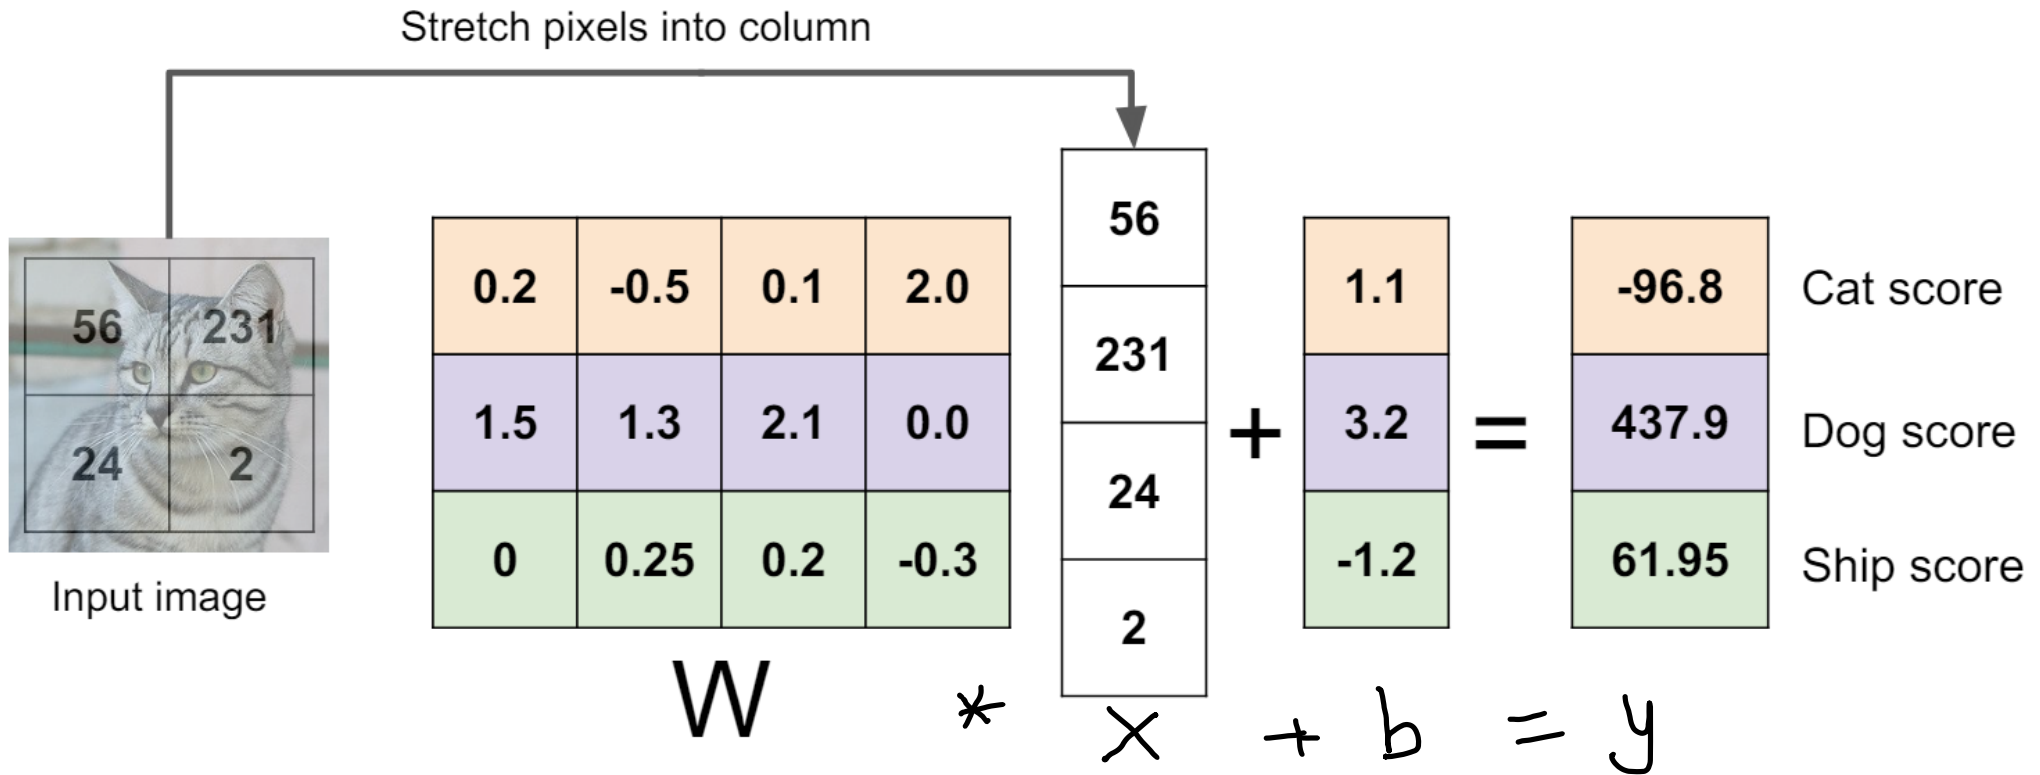


10 numbers,
indicating class
scores



The class with the highest score
is the classification result.

Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



LOSS function



$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

unnormalized probabilities

cat

3.2

car

5.1

frog

-1.7

exp

24.5

164.0

0.18

normalize

0.13

0.87


0.00

$$L_i = -\log(0.13) = 0.89$$

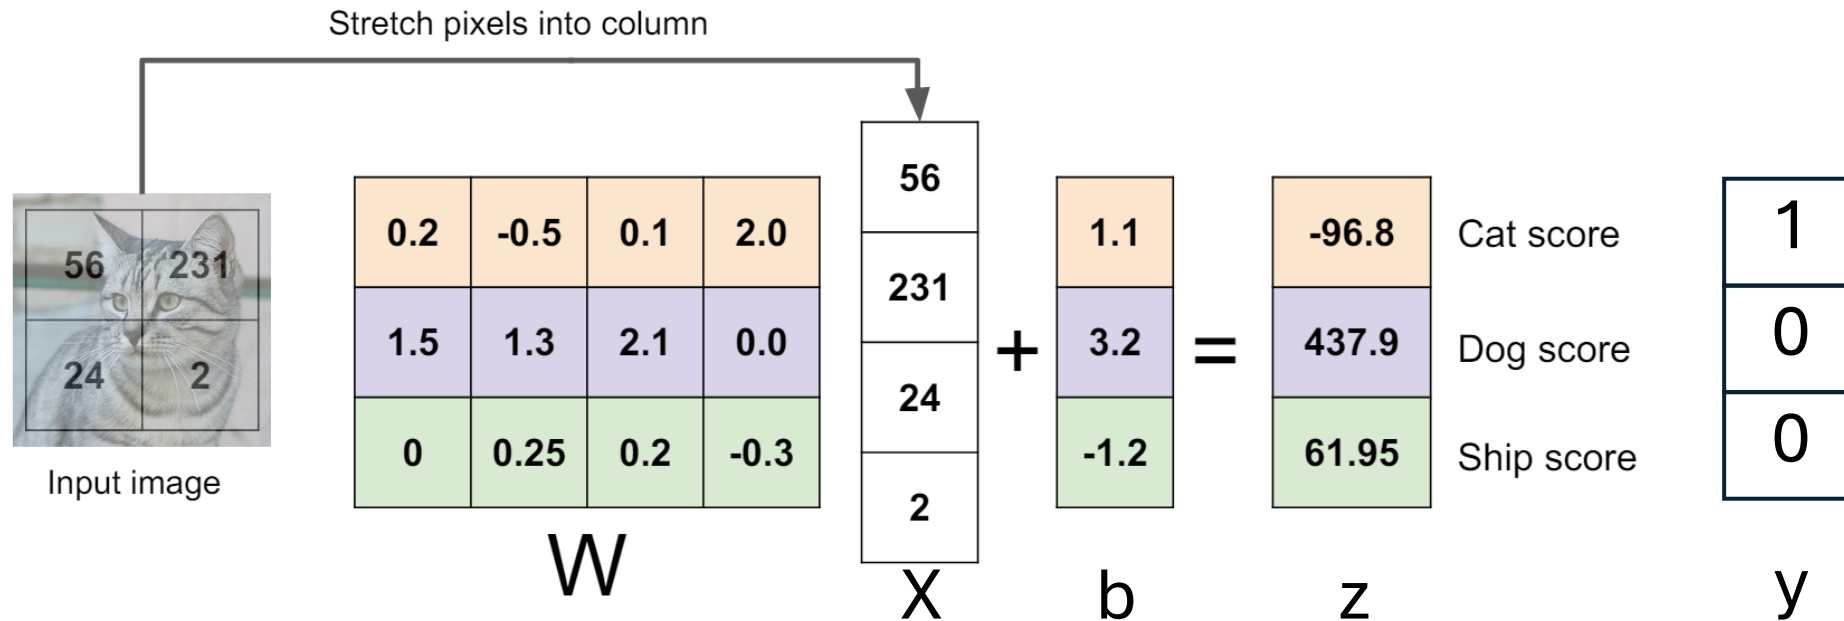
unnormalized log probabilities

probabilities


Minimize the loss



Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



1. Given a few images for training
2. Randomly initialize the weight (W) and offset (b) matrix
3. Calculate z by $z=Wx+b$
4. Compare z and y , calculate loss L
5. Calculate gradients dL/dw , dL/db
6. Update the weight ($W2=W1-\alpha dL/dw$), and offset ($b2=b1-\alpha dL/db$)
7. Iterate from 3 to 6 until the loss becomes convergent



App 1, regression

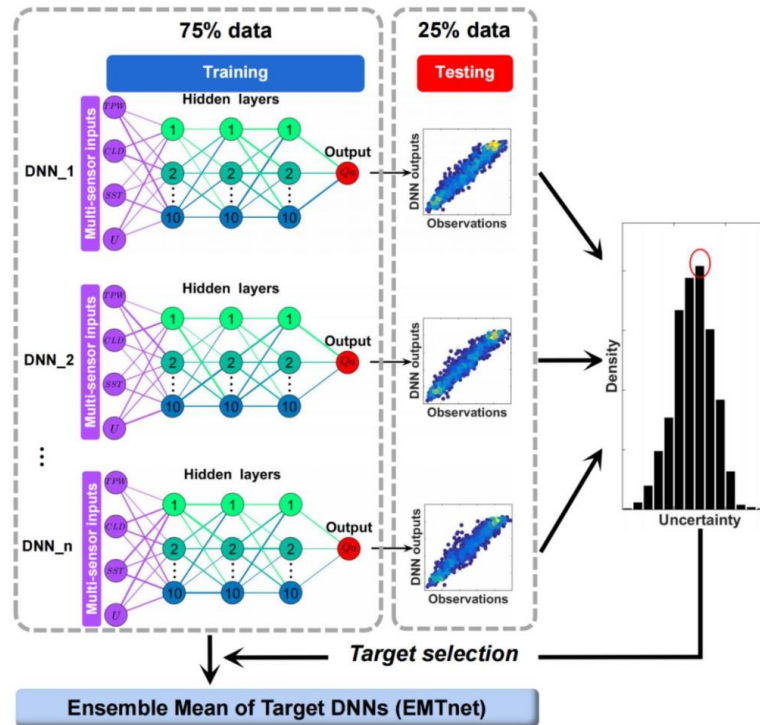


Figure 3. Architecture of the EMTnet. Each DNN_{*n*} ($n = 1, 2, 3, \dots$) uses 75% randomly sampled data from all in situ observations as training data, while the remaining 25% are used as testing data.

Cloud liquid water	U	SST	total column precipitable water
12	12	13	132
121	121	121	232
1212	121	121	34
...

$Wx+b$
→
Near-surface humidity

0

0

0

0

W

*

12

12

13

132

X

+

0

b

=

0

z

32

y

0

0

...

Z=

32

33

...

y=

Loss

$J = \frac{1}{n} \sum_{i=1}^n L_i$

where $L_i = (y_i - z_i)^2$

```
def initialize_parameters(width,Input_W=None,Input_b=None):
```

```
    if Input_W == None:
        Input_W=np.zeros((width,1))
    if Input_b == None:
        Input_b=0
    return Input_W,Input_b
```

```
def model_forward(Input_x,Input_W,Input_b):
```

```
    length,width=np.shape(Input_x)
    z=np.zeros((length,1))
    for i in range(length):
        z[i]=np.dot(Input_W.T,Input_x[i,:])+Input_b
    return z
```

```
def compute_cost(Input_y,Input_z):
```

```
    length=len(Input_y)
    L=np.zeros((length,1))
    for i in range(length):
        L[i]=(Input_y[i]-Input_z[i])**2
    J=np.mean(L)
    return J
```

```
def model_backward(Input_x, Input_y, Input_z):
```

```
    length,width=np.shape(Input_x)
    # calculate dJ/db
    temp1=np.zeros((length,1))
    for i in range(length):
        temp1[i]=(1/length)*(2*Input_z[i]-2*Input_y[i])
    dJdb=np.sum(temp1)

    #calculate dJ/dwj
    dJdwj=np.zeros((width,1))
    for j in range(width):
        temp2=np.zeros((length,1))
        for i in range(length):
            temp2[i]=Input_x[i,j]*(2*Input_z[i]-2*Input_y[i])/length
        dJdwj[j]=np.sum(temp2)

    return dJdb, dJdwj
```

```
def update_parameters(Input_W,Input_b,Input_dJdb,Input_dJdwj,Input_learnrate):
```

```
    width=len(Input_W)
    new_W=np.zeros((width,1))
    for j in range(width):
        new_W[j]=Input_W[j]-Input_learnrate*Input_dJdwj[j]
    new_b=Input_b-Input_learnrate*Input_dJdb

    return new_W,new_b
```

```
def predict(Input_x,Input_W,Input_b):
```



```
    m,n=np.shape(Input_x)
    z=np.zeros((m,1))
    for i in range(m):
        z[i]=np.dot(Input_W.T,Input_x[i,:])+Input_b
    return z
```

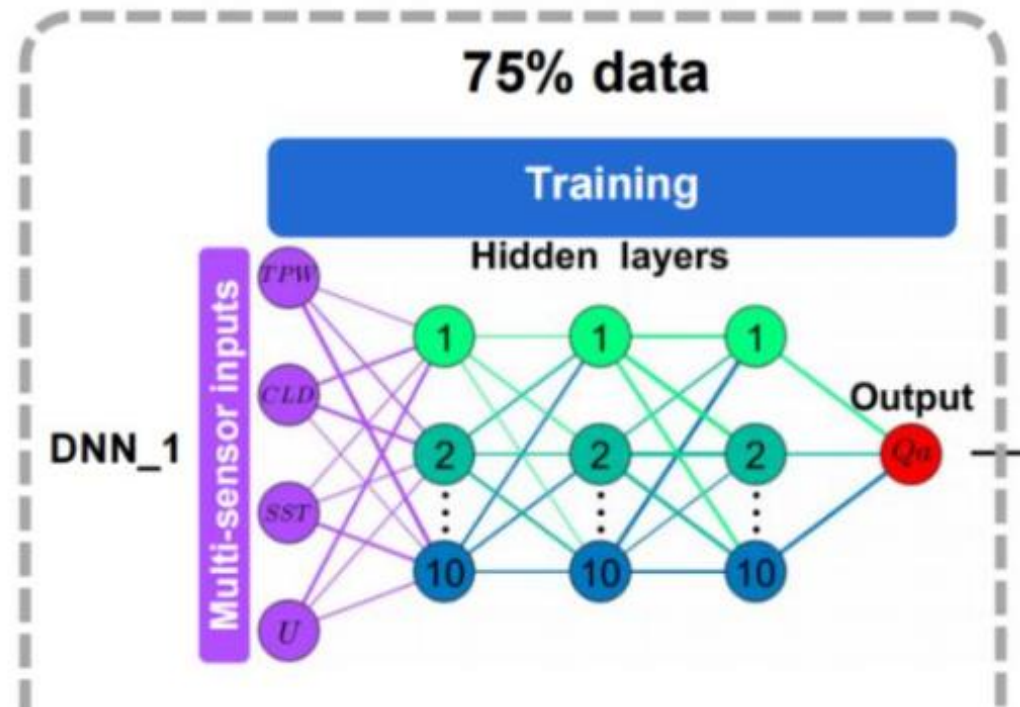

```
def train_linear_model(x,y,iteration,learnrate):
```


```
    m,n=x.shape
    W,b=initialize_parameters(width=n)
    J=np.zeros((iteration,1))

    for itera in range(iteration):
        print('Completed '+str(100*itera/iteration)+'% of work...')
        z=model_forward( x, W, b)
        J[itera]=compute_cost(y, z)
        dJdb,dJdw = model_backward(x, y, z)
        W,b=update_parameters(W, b, dJdb, dJdw, learnrate)

    return W,b,J
```





$h(g(f(x)))$

1st layer:

0	0	0	0
0	0	0	0
0	0	0	0

W1

12
12
13
132

X

*

0
0
0

b1

+

=

0
0
0

z1

2st layer:

0	0	0
---	---	---

W2

*

0
0
0

z1

+

0

b2


=

0


z2

Qa


32



1 linear layer




2 linear layer




3 linear layer


Activation functions



$$\text{sigmoid}(x) = \frac{1}{1+e^{-x}}$$

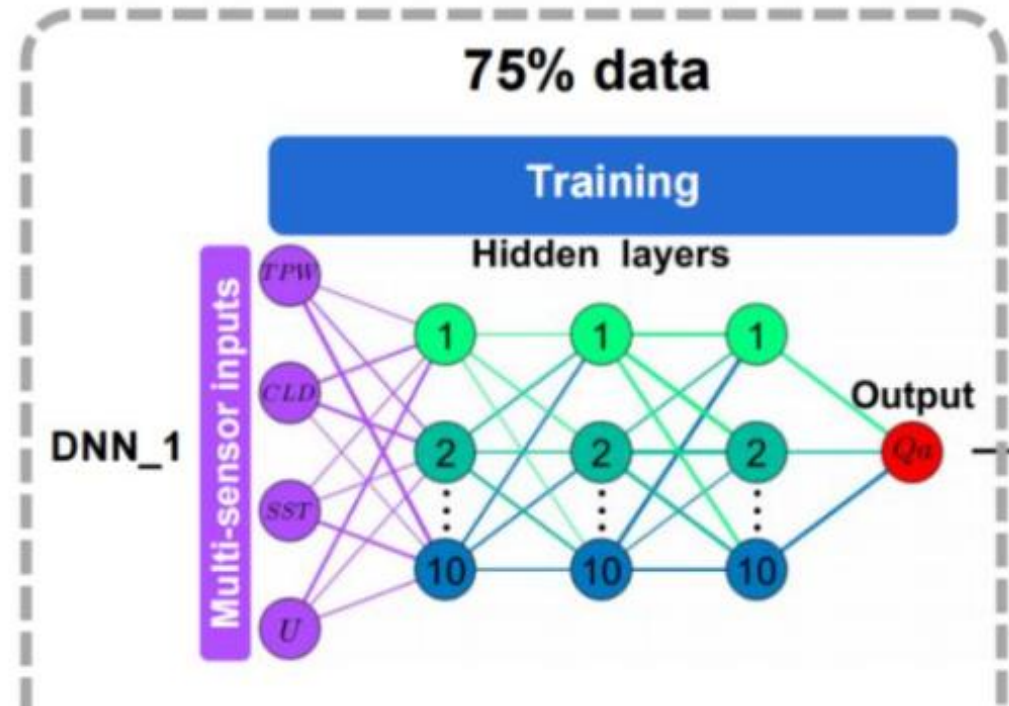


$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} = 2 \text{sigmoid}(2x) - 1$$





$$\text{ReLU}(x) = \max(0, x)$$

“Rectified Linear Unit” (diode function)



Hidden Layers, Black Box, Architecture, neural network
Artificially designed,
No scientific interpretation

Neuron and Artificial neuron



Src: <http://www.theprojectspot.com/tutorial-post/introduction-to-artificial-neural-networks-part-1/7>

$$output = h\left(\sum_{i=1}^n x_i w_i\right)$$

App2: ML for weather prediction




Figure 2. Schematic illustrating the architecture of our DLWP CNN based on the U-Net architecture. Each red arrow represents a 2-D convolution operating on each cubed-sphere face. Green and purple arrows indicate average-pooling and upsampling operations, respectively. The blue-to-yellow lines represent skip connections, whereby the blue state is copied exactly to the yellow state vector and concatenated to the new blue state vector along the channels dimension. The final gray arrow is a 1×1 convolution. The blue numbers indicate the number of convolutional filters (channels) at each stage of the network (channel width is to scale).

2D data: CNN (convolutional neural network)







Figure 2. MetNet Structure.

Fully connected layers

vs


convolutional neural network





Alexnet

2012



Imagenet classification with deep convolutional neural networks

[A Krizhevsky, I Sutskever...](#) - Advances in neural ..., 2012 - proceedings.neurips.cc

We trained a large, deep convolutional neural network to classify the 1.3 million high-resolution images in the LSVRC-2010 ImageNet training set into the 1000 different classes. On the test data, we achieved top-1 and top-5 error rates of 39.7\% and 18.9\% which is considerably better than the previous state-of-the-art results. The neural network, which has 60 million parameters and 500,000 neurons, consists of five convolutional layers, some of which are followed by max-pooling layers, and two globally connected layers with a final ...

☆ Save 77 Cite Cited by 128198 Related articles All 102 versions 77

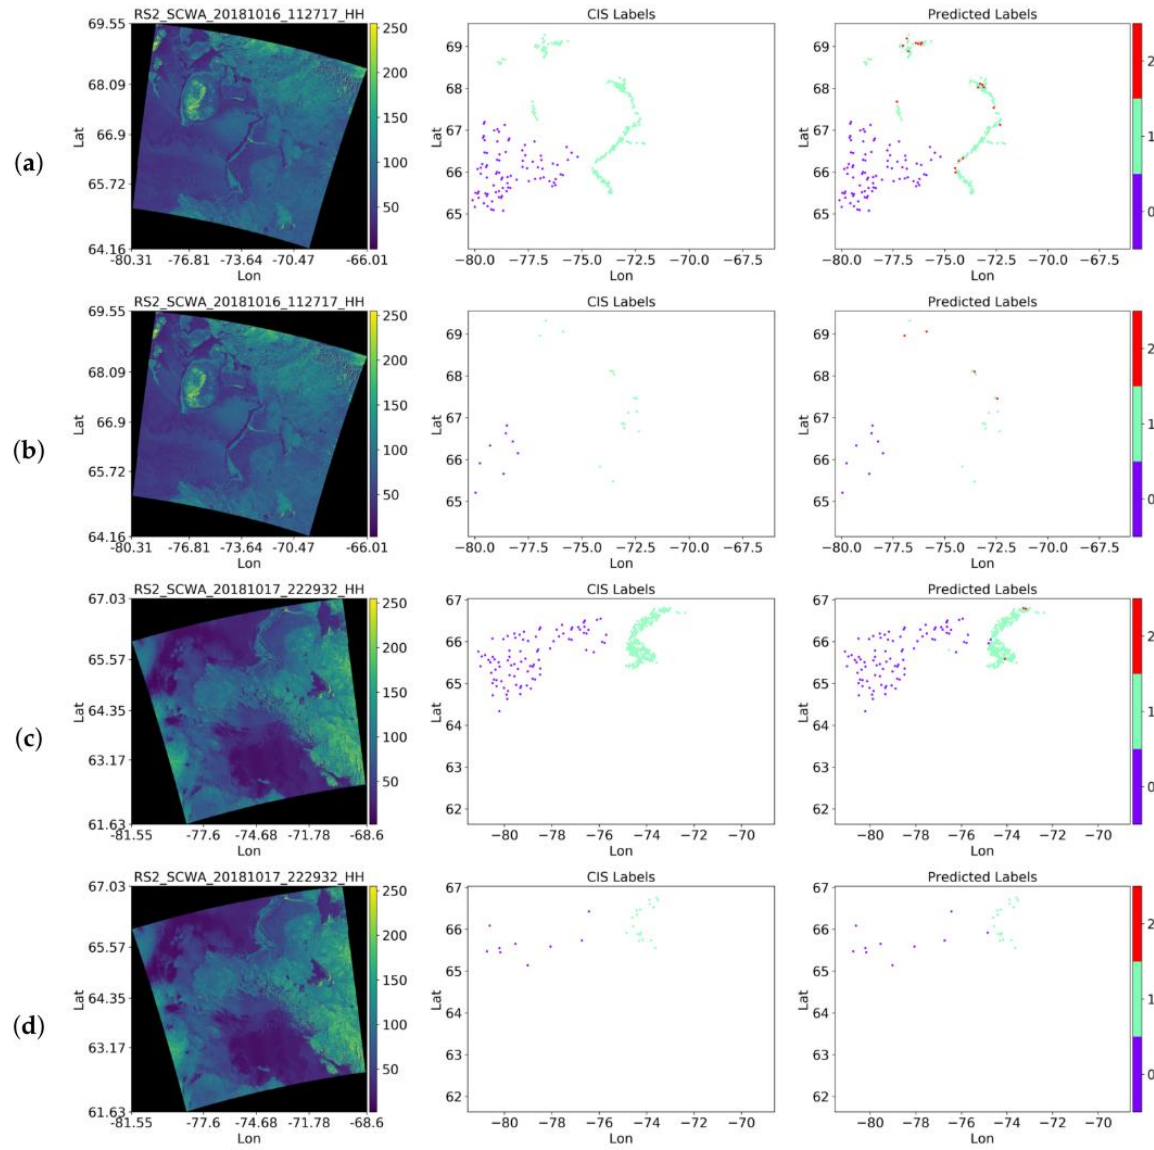
Import pytorch

Use bricks in *pytorch* to build your own architecture.

```
# Define the CNN model
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        # First convolutional layer
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3, padding=1)
        # Second convolutional layer
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        # Max pooling layer
        self.pool = nn.MaxPool2d(2, 2)
        # First fully connected layer
        self.fc1 = nn.Linear(64 * 7 * 7, 128)
        # Second fully connected layer
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        # First convolutional layer followed by ReLU activation and pooling
        x = self.pool(torch.relu(self.conv1(x)))
        # Second convolutional layer followed by ReLU activation and pooling
        x = self.pool(torch.relu(self.conv2(x)))
        # Flatten the tensor
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        # First fully connected layer followed by ReLU activation
        x = torch.relu(self.fc1(x))
        # Second fully connected layer
        x = self.fc2(x)
        return x
```


App 3, image segmentation



Proof of Concept for Sea Ice Stage of Development
Classification Using Deep Learning



Input:
 $3 \times H \times W$



Predictions:
 $H \times W$

0	1	0	1	1
0	0	0	1	1
0	0	0	1	1
2	0	0	3	3
2	2	2	2	3

1: sky


0: cow

2: grass

3: tree

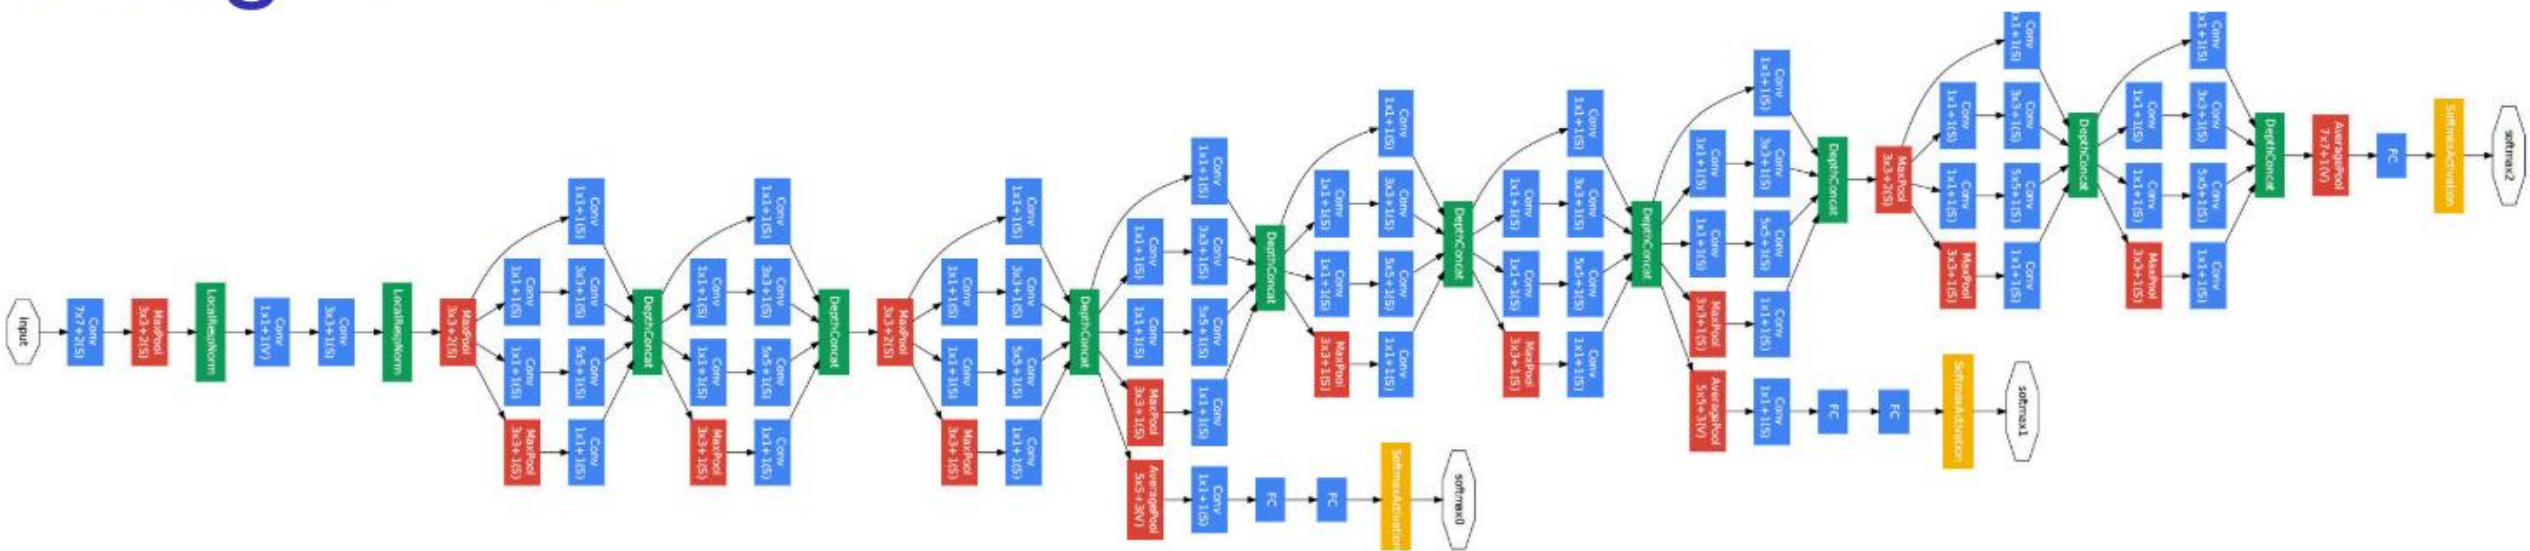
Outlook: deeper layers

Deeeeep




Outlook: wider layers

GoogLeNet




Improving machines → teaching humans

1. When AI is weaker than humans
 - Transparency → finding error modes
 - Goal = improving machines
2. When AI is at par with humans
 - Transparency → providing rationales
 - Goal = building trust with humans to drive adoption
3. When AI is stronger than humans
 - Transparency → explaining a complicated concept
 - Goal = teaching humans



- Let's take a single value in an intermediate feature map and propagate its gradient back to the original image pixels
- What does this tell us?



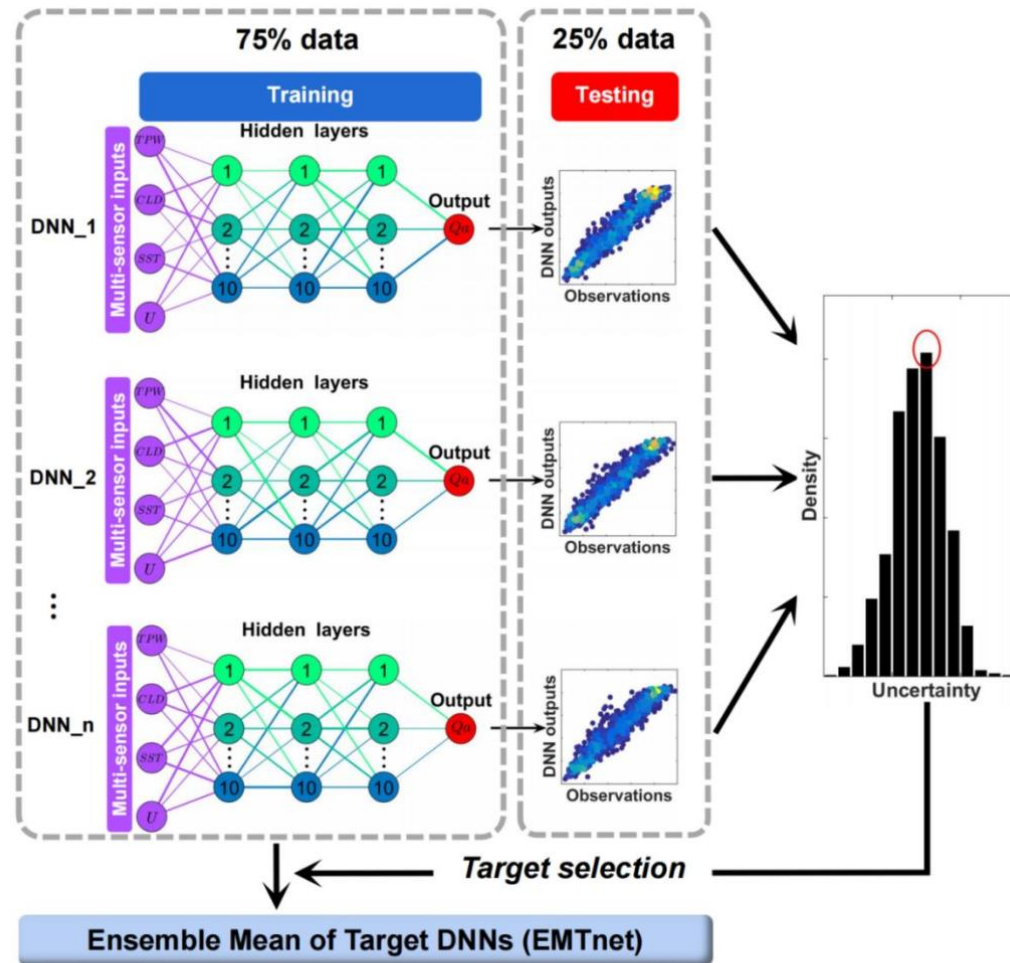
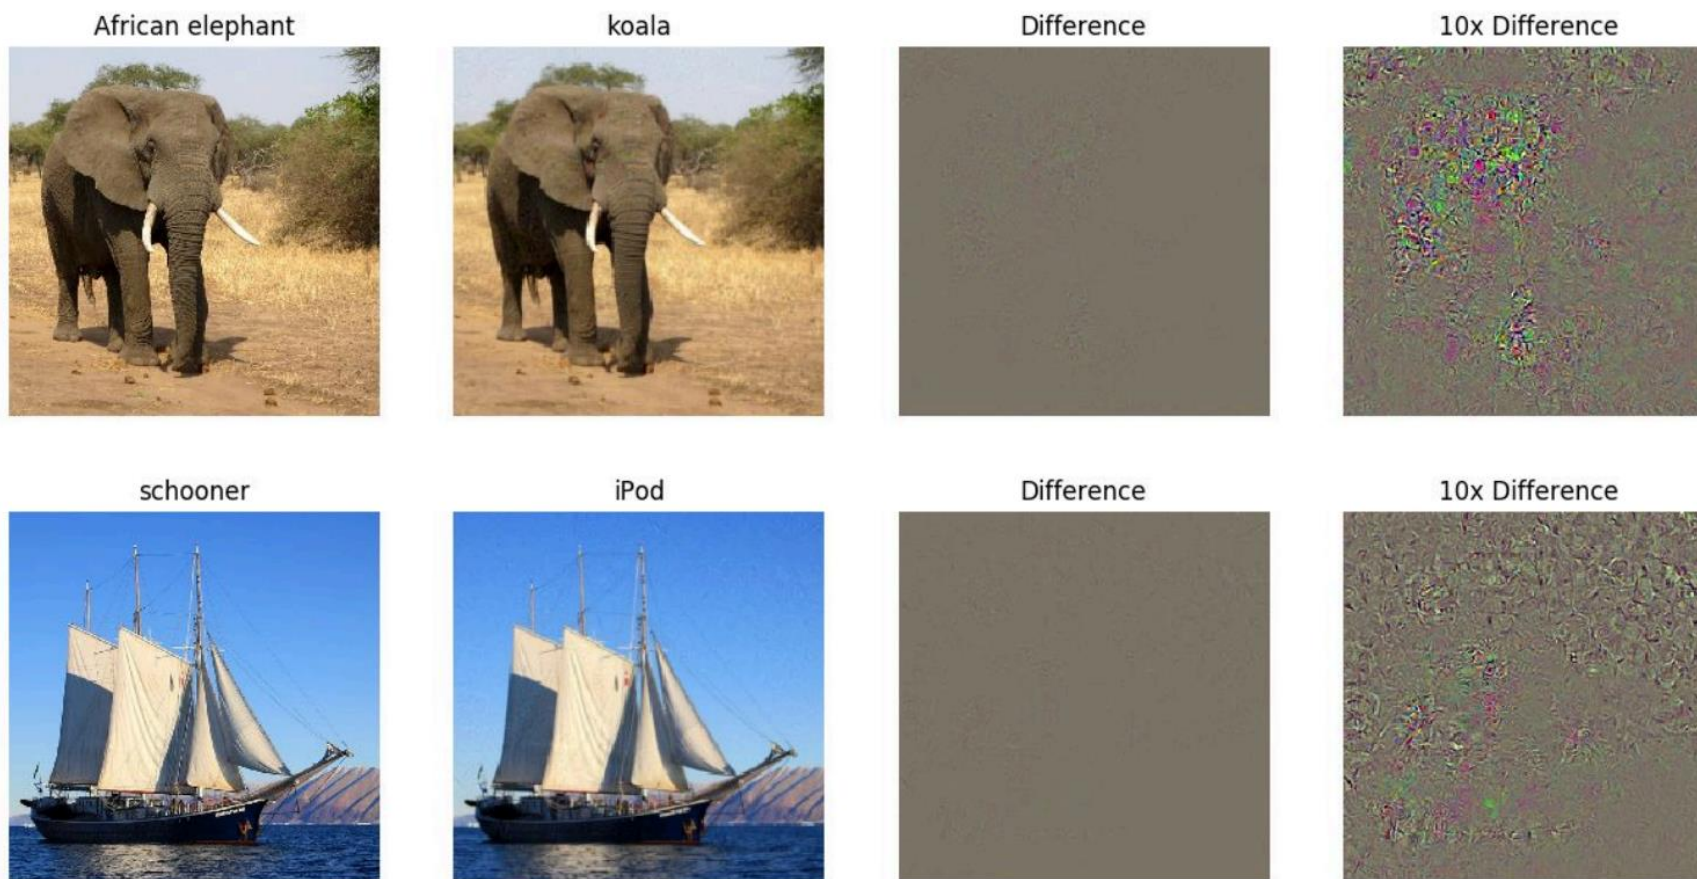
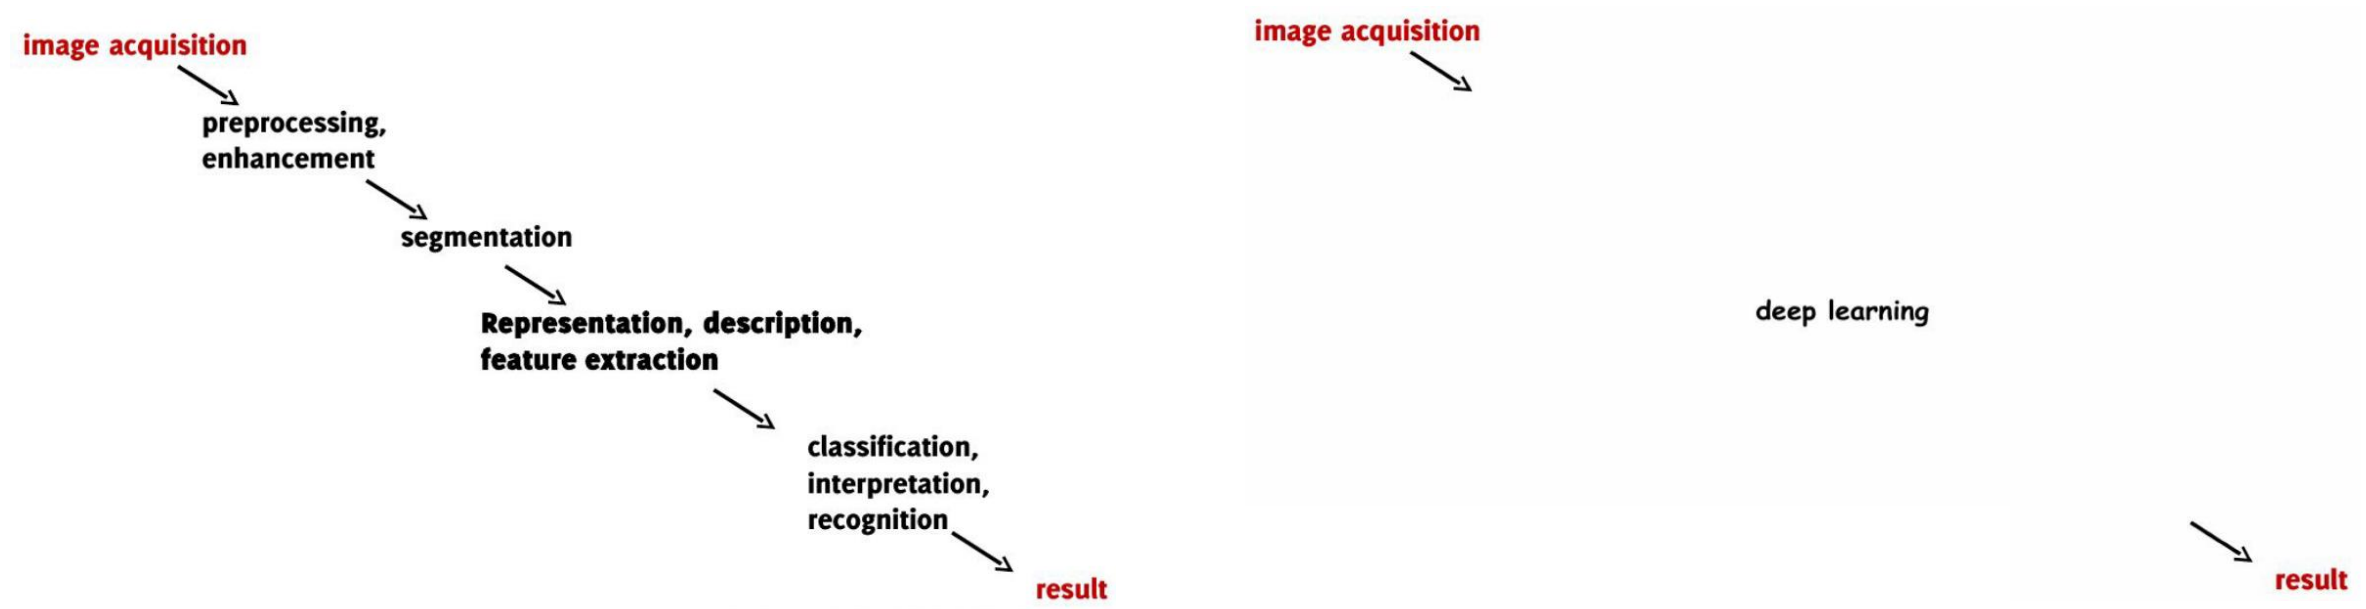


Figure 3. Architecture of the EMTnet. Each DNN_{*n*} (*n* = 1, 2, 3, ...) uses 75% randomly sampled data from all in situ observations as training data, while the remaining 25% are used as testing data.

Trained model is still Fragile



Conclusion



Conventional method VS machine learning

Conclusion

- Machine Learning is not always the best tool, and need to be used with care, augmentation, texture, bias...
- Can be rather demanding – lots of annotated data, compute heavy
- ML is great at interpolation, but less great at extrapolation
- Don't use it to solve known maths:
 - > ML is more useful to provide a starting guess
 - > In many cases, combinations of classic and learning based methods provide the best solutions